

NEURAL COMPUTATION

Prof. Alexander Huth

3.8.2021

RECAP

$$Y = f(X)$$

- * System identification
 - * Linear
 - * Linearized
 - * Nonlinear

RECAP

$$Y = f(X)$$

- * System identification

- * Linear

$$Y = X\beta$$

- * Linearized

$$Y = \mathbb{L}(X)\beta$$

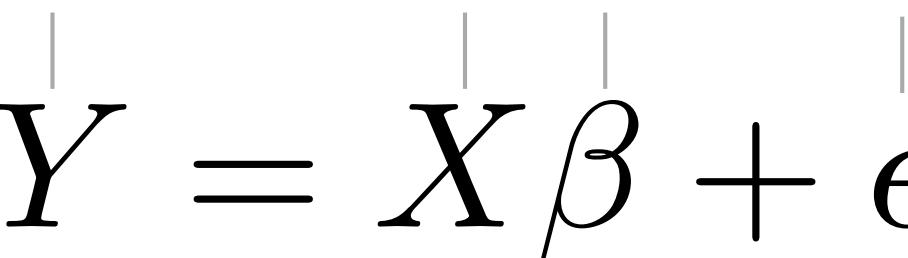
- * Nonlinear

$$Y = \Theta(X)$$

LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



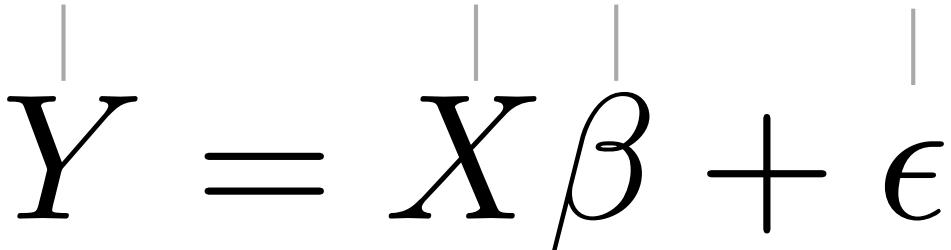
Loss function:

$$\|Y - X\beta\|_2$$

LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE

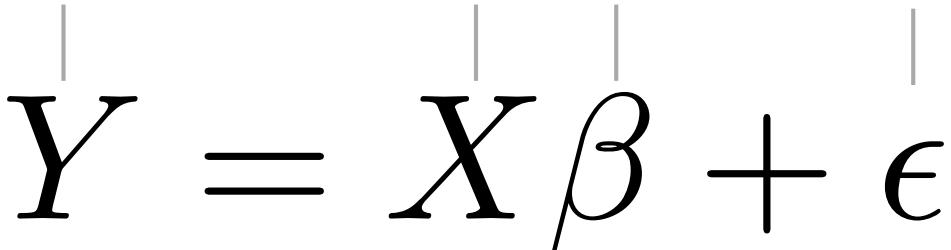


- * How do we solve for beta?
 - * Analytically
 - * Iteratively

ANALYTIC SOLUTION TO LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



$$\beta = f(X, Y)$$

ITERATIVE SOLUTION TO LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



$$\Delta\beta \propto -\frac{\partial(||Y - X\beta||_2)}{\partial\beta}$$

LINEAR REGRESSION

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



Constraining the values **beta** can take improves model performance:

REGULARIZATION

REGULARIZATION

- * Regularization can be thought of in three ways:
 - * **Prior**
 - * **Penalty**
 - * **Geometry**

REGULARIZATION AS PRIOR

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



$$Y_{t,j} \sim \mathcal{N}(X\beta, \sigma^2)$$

$$\hat{\beta} = \operatorname*{argmax}_{\beta} P(Y|X, \beta)$$

REGULARIZATION AS PRIOR

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



$$Y_{t,j} \sim \mathcal{N}(X\beta, \sigma^2)$$

$$\hat{\beta} = \operatorname{argmax}_{\beta} P(Y|X, \beta)P(\beta)$$

REGULARIZATION AS PENALTY

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE



$$E(\beta) = \|Y - X\beta\|_2^2$$

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} E(\beta)$$

REGULARIZATION AS PENALTY

$$Y = X\beta + \epsilon$$

RESPONSE VARIABLES WEIGHTS NOISE

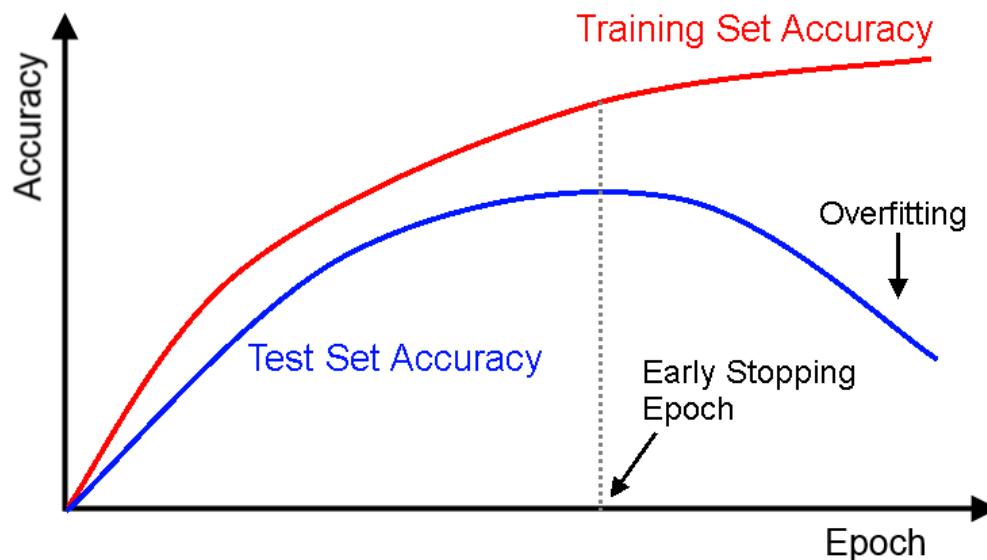


$$E_{pen}(\beta) = ||Y - X\beta||_2^2 + \lambda ||\beta||_2^2$$

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} E_{pen}(\beta)$$

REGULARIZATION AS GEOMETRY - EARLY STOPPING

$$\begin{array}{cccc} \text{RESPONSE} & \text{VARIABLES} & \text{WEIGHTS} & \text{NOISE} \\ | & | & | & | \\ Y & = X\beta + \epsilon \end{array}$$



COMMON TYPES OF REGULARIZATION

- * **Beta is small (L2-sense)** = ridge = gradient descent w/ early stopping
- * **Beta is small, sparse (L1-sense)** = LASSO = coord. descent w/ early stopping
 - * Beta is small & sparse (L1+L2 sense) = elastic net
- * **Beta is sparse (L0-sense)** = variable selection

RIDGE REGRESSION

- * Multivariate normal (MVN) prior on beta
- * L2 penalty on beta
- * Gradient descent w/ early stopping

RIDGE REGRESSION

$$Y = X\beta + \epsilon$$

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} [||Y - X\beta||_2^2 + \lambda ||\beta||_2^2]$$

ERROR or LOSS **PENALTY**

RIDGE REGRESSION

$$\hat{\beta} = (X^\top X + \lambda I)^{-1} X^\top Y$$

$$\hat{\beta} = X_{ridge}^+ Y$$

RIDGE REGRESSION

- * Efficient solution with SVD

$$\hat{\beta} = (X^\top X + \lambda I)^{-1} X^\top Y$$

$$(\text{SVD}) \quad X = U S V^\top \quad D = \frac{S}{S^2 + \lambda^2}$$

$$\hat{\beta} = V D U^\top Y$$

RIDGE REGRESSION

- * How to choose lambda?
- * GCV - Generalized Cross Validation 🤔
- * Block-wise cross-validation 👍

RIDGE REGRESSION

- * Good implementation: scikit-learn
- * Awesome implementation:
<http://github.com/alexhuth/ridge>

LASSO

- * Laplacian prior on β_i
- * L1 penalty on β
- * Coordinate descent w/ early stopping

LASSO

$$Y = X\beta + \epsilon$$

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} [||Y - X\beta||_2^2 + \lambda ||\beta||_1]$$

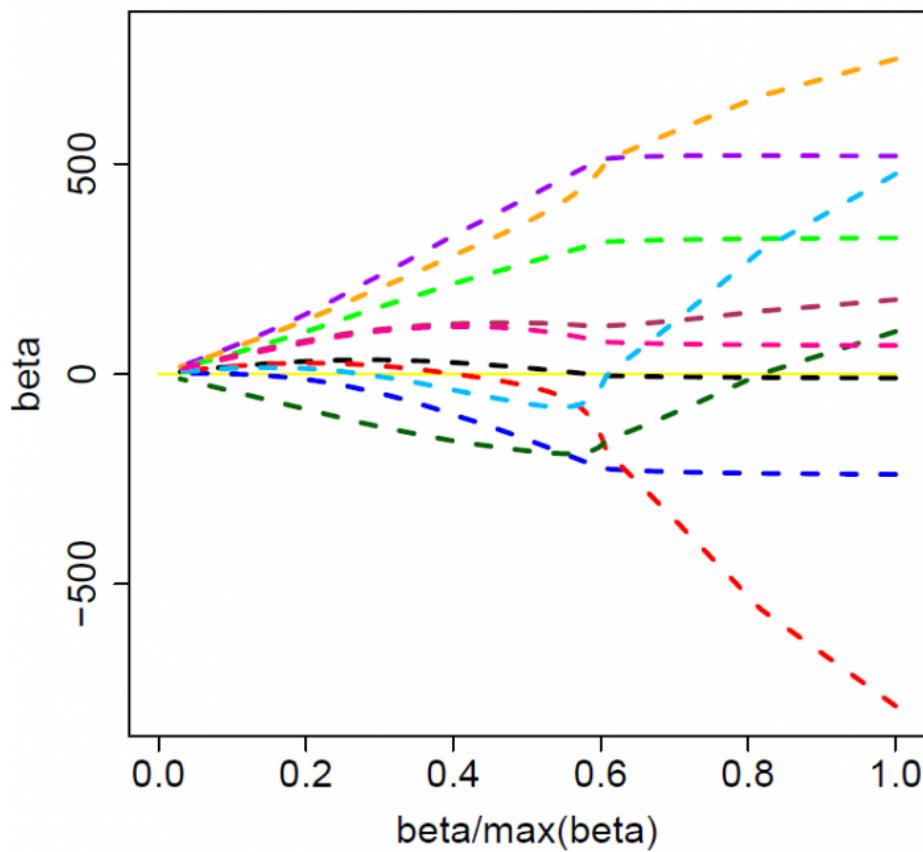
ERROR or LOSS PENALTY

LASSO

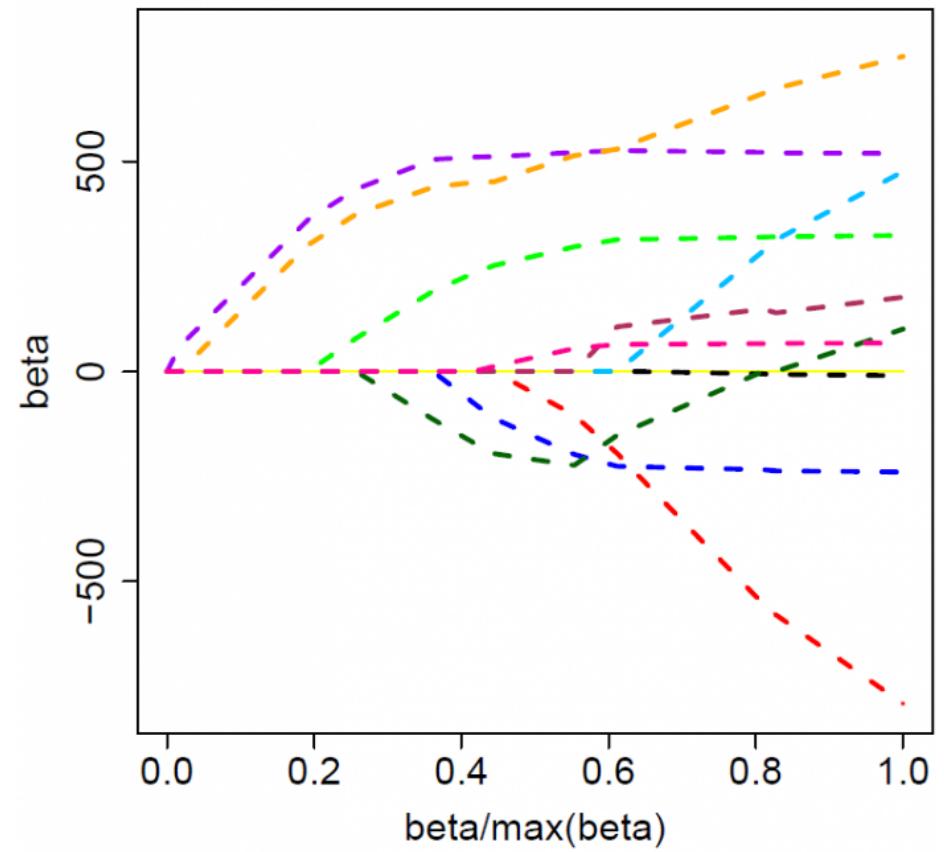
- * No closed form solution
- * Solved via coordinate descent, LARS (least-angle regression) or other methods
- * *SssssLLLlooooowWWWW.....*

LASSO

Ridge Regression



Ordinary Lasso



OTHER METHODS

- * Neural networks (linear or nonlinear)
- * Random forests (nonlinear)
- * Feature selection ($\sim L_0$ -norm)

NEXT TIME

- * Tikhonov regression